

# Widespread Error Detection in Large Scale Continuous Integration Systems

---

- Stanislaw Swierc (stansw)
- James Lu (luchangj)
- Thomas Yi (thomasyi)

\* @meta.com

**<https://github.com/StanislawSwierc/CCIW2024-Widespread-Error-Detection>**



# Disclaimer

*This presentation contains sample code with a basic pipeline for detecting widespread errors in the React open-source project. It does NOT contain any code, logs or data internal to Meta Platforms, Inc.*

# Motivation

*Improve developer productivity by quickly detecting and mitigating the impact of widespread errors.*

## Motivation: Widespread Errors

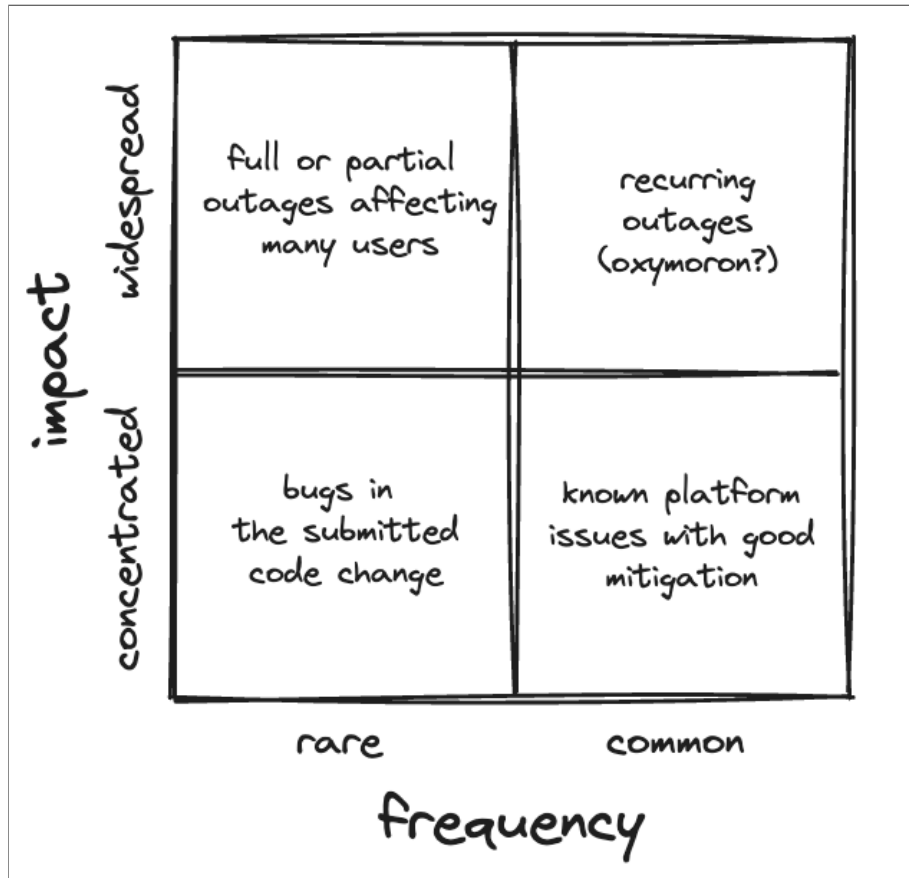
Widespread errors are commonly infrastructural errors that falsely block developers from integrating their code changes into the mainline branch.

Why is this bad?

- **Productivity** - Developers waste time trying to fix code, when it is not at fault.
- **Trust** - Erodes confidence in CI System's ability to produce validation signals.
- **Efficiency** - Need to re-run failing validation, thus, wasting machine resources.

## Motivation: Impact vs. Frequency

We can organize errors by number of affected people and how common they are:



### Challenges

- Rare, widespread errors may appear like a group of concentrated errors.

- Capturing signatures of rare, widespread errors is only good for retroactive analysis, but not for detection.

# React Case Study

React uses CircleCI to validate pull requests and the health of the mainline branch.

- **Project** - validation configuration for the code repository
  - **Pipeline** - group of workflows
    - **Workflow** - graph of jobs
      - **Job** - set of steps executed sequentially
        - **Step/Action** - execution unit
          - **Logs** - standard output and summary

<https://circleci.com/docs/concepts>

# React Case Study: CircleCI API

CircleCI offers an API to fetch information about pipelines and other resources all the way to individual logs.

In [3]:

```
from pycircleci.api import Api

ORG = "facebook"
PROJECT = "react"
TOKEN = "{token}"

client = Api(token=TOKEN)
```

In [37]:

```
%%script true

client.get_project_pipelines(...)
client.get_pipeline_workflow(...)
client.get_workflow_jobs(...)
client.get_job_details_v1(...)
client._session.get(action["output_url"])
```

In [30]:

```
!du -h pipelines.jsonl
```

```
4.4G    pipelines.jsonl
```



**<https://circleci.com/docs/api/v2/index.html>**

**<https://circleci.com/docs/managing-api-tokens>**

## # React Case Study: Actions Fact Table

In [29]:

```
actions.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 341630 entries, 0 to 341629
```

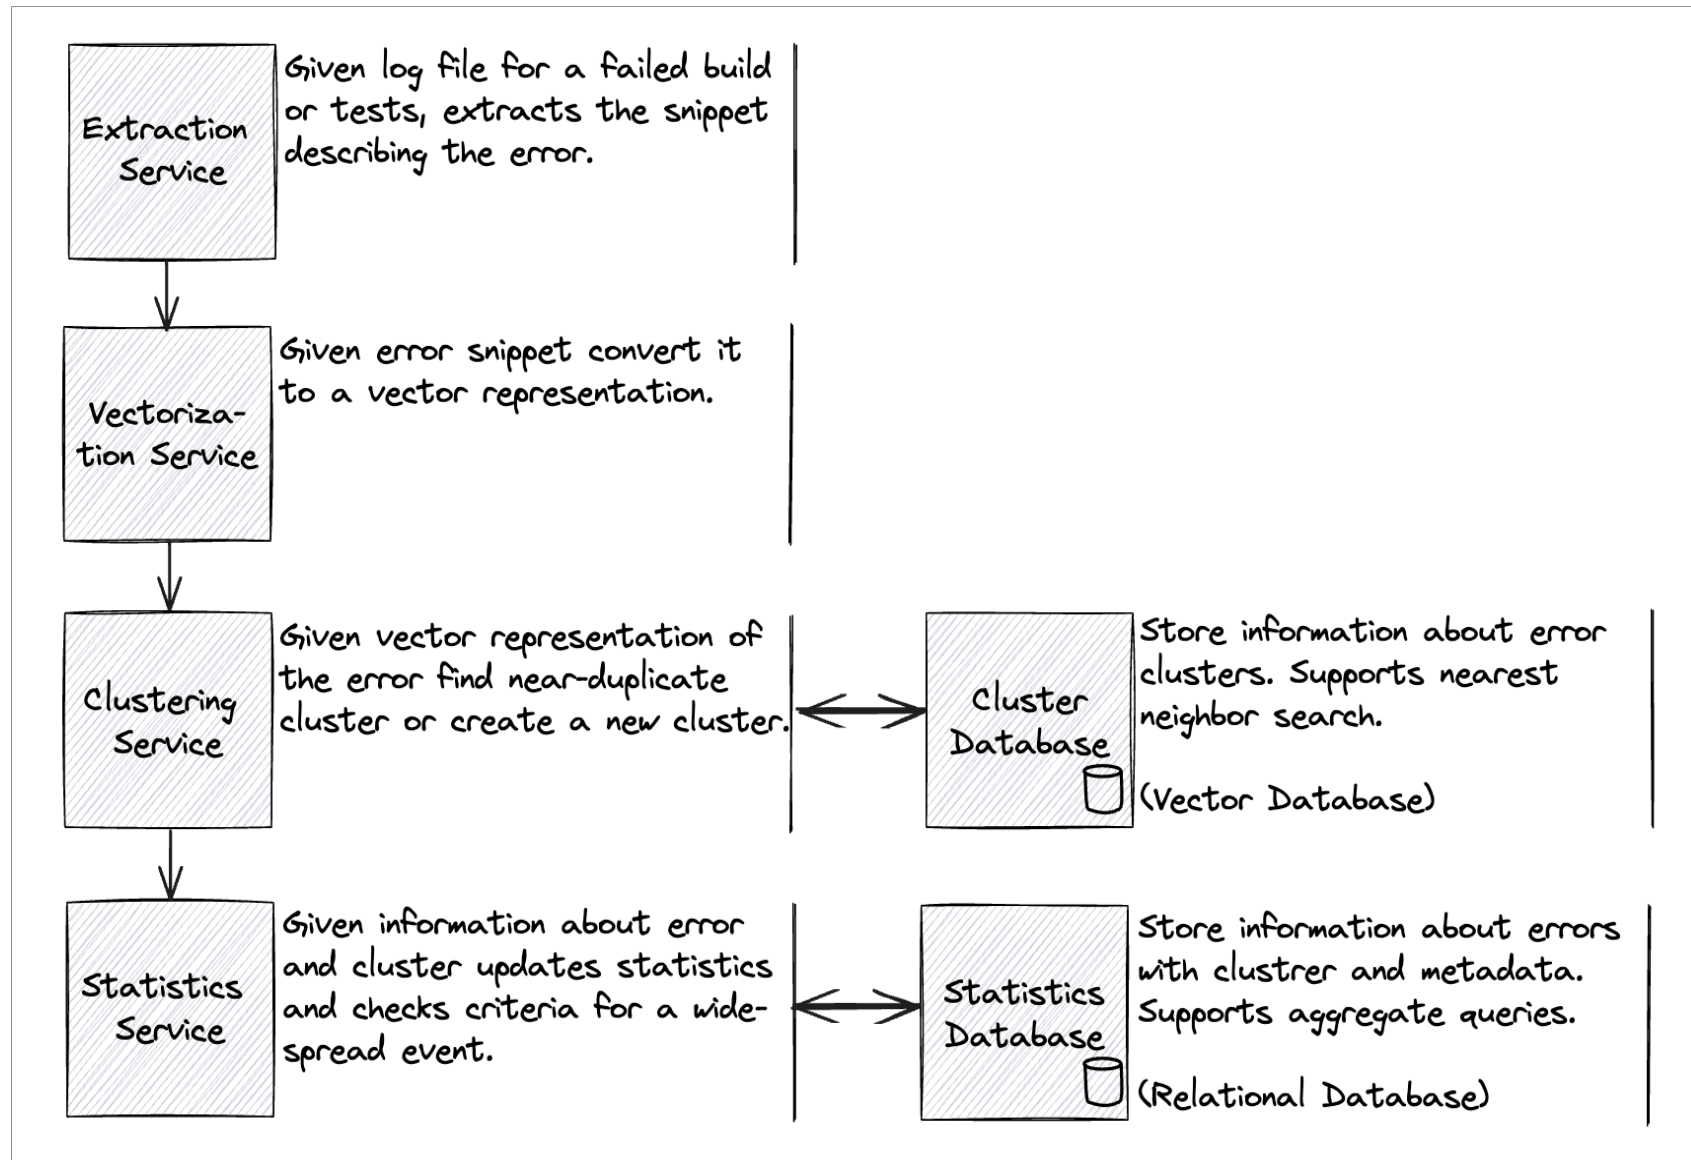
```
Data columns (total 17 columns):
```

#	Column	Non-Null Count	Dtype
0	pipeline_id	341630 non-null	object
1	pipeline_actor_login	341630 non-null	object
2	workflow_id	341630 non-null	object
3	workflow_name	341630 non-null	object
4	workflow_status	341630 non-null	object
5	job_id	341630 non-null	object
6	job_name	341630 non-null	object
7	job_status	341630 non-null	object
8	job_link	341630 non-null	object
9	action_name	341630 non-null	object
10	action_fully_qualified_name	341630 non-null	object
11	action_index	341630 non-null	int64
12	action_status	341630 non-null	object
13	action_output_time	341630 non-null	datetime64[ns, UTC]
14	action_output_message	341630 non-null	object
15	action_output_type	341630 non-null	object
16	action_output_date	341630 non-null	datetime64[ns, UTC]

```
dtypes: datetime64[ns, UTC](2), int64(1), object(14)
```

```
memory usage: 44.3+ MB
```

# Pipeline Design



Extraction

# Extraction

Given a log file for a failed build or tests, extracts the snippet describing the error.

1. **Federated solution** - allow teams to provide custom rules for extracting relevant error snippets.
2. **General solution** - find lines which appear only in logs of failed actions.

# Extraction: Example

✖ yarn test --build --project=devtools -r=experimental --ci

✦ Explain this error AI Experiment 10s 🔍 📄 ⬇

```
#!/bin/bash -eo pipefail
yarn test --build --project=devtools -r=experimental --ci

1 yarn run v1.22.5
2 $ node ./scripts/jest/jest-cli.js --build --project=devtools -r=experimental --ci
3 $ NODE_ENV=development RELEASE_CHANNEL=experimental compactConsole=false node ./scripts/jest/jest.js --config ./s
  cripts/jest/config.build-devtools.js --maxWorkers=2
4
5 Running tests for devtools (experimental)...
6 FAIL packages/react-devtools-shared/src/__tests__/utils-test.js
7   • Test suite failed to run
8
9   Cannot find module 'react/src/ReactElement' from 'packages/react-devtools-shared/src/__tests__/utils-test.js'
10
11     24 |   REACT_STRICT_MODE_TYPE as StrictMode,
12     25 | } from 'shared/ReactSymbols';
13   > 26 | import {createElement} from 'react/src/ReactElement';
14     |                                     ^
15     27 |
16     28 | describe('utils', () => {
17     29 |   describe('getDisplayName', () => {
18
19     at Resolver._throwModNotFoundError (node_modules/jest-resolve/build/resolver.js:427:11)
20     at Object.<anonymous> (packages/react-devtools-shared/src/__tests__/utils-test.js:26:21)
21
22 PASS packages/react-devtools-shared/src/__tests__/legacy/inspectElement-test.js (6.106 s)
23
24 Test Suites: 1 failed, 1 passed, 2 total
25 Tests:      13 passed, 13 total
26 Snapshots:  12 passed, 12 total
27 Time:       7.372 s
28 Ran all test suites.
29 error Command failed with exit code 1.
30 info Visit https://yarnpkg.com/en/docs/cli/run for documentation about this command.
31
32 Exited with code exit status 1

CircleCI received exit code 1
```

Runtime information

Failed tests

Passed tests

Summary

## Extraction: Sample Code

In [183]:

```
def assign_action_output_failure(df, cache_size=2048):
    result = []
    action_cache = defaultdict(lambda: cachetools.LRUCache(cache_size))
    for i, row in tqdm(df.iterrows(), total=len(df)):
        failure_lines = []
        success_cache = action_cache[row["action_fully_qualified_name"]]
        for line in row["action_output_message"].splitlines():
            template = parse_line(line)


            if row["action_status"] == "success":
                success_cache[template] = True
            elif not template or template not in success_cache:
                failure_lines.append(line)

        result.append("\n".join(failure_lines))
    return result
```

# Extraction: Results

In [433]:

```
line_num_ecdf_fig.show()
```





## Extraction: Extensions

1. Use log parser and diff failures based on log templates

*Jiang, Z., Liu, J., Huang, J., Li, Y., Huo, Y., Gu, J., Chen, Z., Zhu, J. and Lyu, M.R., 2023.*

*A Large-scale Benchmark for Log Parsing*

# Vectorization

# Vectorization

Convert raw text data into a numerical format that can be efficiently processed.

1. Bag-of-words model with open-vocabulary
2. Minhash (e.g. with 256 x 16-bit we get a constant 0.5 kB per error)

$$\begin{array}{c} \text{Jaccard Distance ( } \boxed{\begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \end{array}} , \boxed{\begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \end{array}} \text{ ) } = d \\ \text{Minhash} \downarrow \qquad \downarrow \\ 1/N * \text{Generalized Hamming Distance ( } [\text{---}] , [\text{---}] \text{ ) } \approx d \end{array}$$

*Leskovec J, Rajaraman A, Ullman JD. "Chapter 3: Finding Similar Items" Mining of massive data sets. Cambridge university press; 2020 Jan 9.*

## Vectorization: Normalization

Reduce the size of the vocabulary by removing common hashes and common metadata.

In [236]:

```
normalize_rules = [  
    ("{{GUID}}", re2.compile(r'[0-9a-f]{8}-[0-9a-f]{4}-[0-9a-f]{4}-[0-9a-f]{4}-[0-9a-f]{12}')),  
]  
  
def normalize(text):  
    for repl, regex in normalize_rules:  
        text = regex.sub(repl, text)  
    return text
```

# Vectorization: Tokenization

Extract meaningful tokens.

In [136]:

```
token_regex = re2.compile(r"\b[[:alpha:]]_[[:alpha:][:digit:]]_+")  
  
def tokenize(text):  
    normalized_text = normalize(text)  
    return token_regex.findall(normalized_text)
```

# Vectorization: Minhash

Convert set of tokens to Minhash sketches.

In [182]:

```
from datasketch import MinHash

NUM_PERM = 256

def vectorize(tokens):
    m = MinHash(num_perm=NUM_PERM)
    for token in set(tokens):
        m.update(token.encode('utf8'))
    return m.hashvalues
```

## Vectorization: Results

Minhash representation is more compact and allows for fast nearest-neighbors search.

Total size of original text data:

In [238]:

```
print("{:.3f} GB".format(actions["action_output_message"].str.len().sum() / 2 ** 30))
```

1.978 GB

Total size of minhash sketches:

In [234]:

```
print("{:.3f} GB".format(actions.shape[0] * np.uint16(0).nbytes * NUM_PERM / 2 ** 30))
```

0.163 GB

## Vectorization: Extensions

1. **Process documents in parallel.** Vectorization is embarrassingly parallelizable.
2. **Use Weighted Minhash.** Minhash with Inverse Document Frequency (IDF) weights can automatically ignore common terms such as "failed".

*Chum, O., Philbin, J. and Zisserman, A., 2008, September. Near duplicate image detection: Min-hash and TF-IDF weighting.*

3. **Use Neural Network Vector Embeddings (future).** Embeddings are a good alternative to Minhash sketches and they can capture semantics of logs.

*Meng, W., Liu, Y., Huang, Y., Zhang, S., Zaiter, F., Chen, B. and Pei, D., 2020, August. A semantic-aware representation framework for online log analysis.*



# Clustering

# Clustering

Given vector representation of the error, find a **near-duplicate cluster** or create a new cluster.

- Near-duplicate search is performed on a vector database.
- For Minhash the database should support *Generalized Hamming Distance*.
- LSH Forest is an efficient data structure and offers time complexity of  $O(n \log(n))$ .

*Bawa, M., Condie, T. and Ganesan, P., 2005, May. LSH forest: self-tuning indexes for similarity search.*

# Clustering: Sample Code

In [274]:

```
MAX_DISTANCE = 0.25

def assign_action_output_cluster_info(df):
    cluster_info = []
    forest = MinHashLSHOnlineForest(num_perm=NUM_PERM, num_trees=16)
    for index, sketch in tqdm(df["action_output_sketch"].items(), total=len(df)):
        neighbors = forest.query(MinHash(hashvalues=sketch), 10)

        if neighbors:
            distances = [jaccard_distance(sketch, x) for x in df.loc[neighbors, "action_output_sketch"]]
            nearest_i = np.argmin(distances)
            nearest_index, nearest_dist = (neighbors[nearest_i], distances[nearest_i])
        else:
            nearest_index, nearest_dist = (0, 1.0)

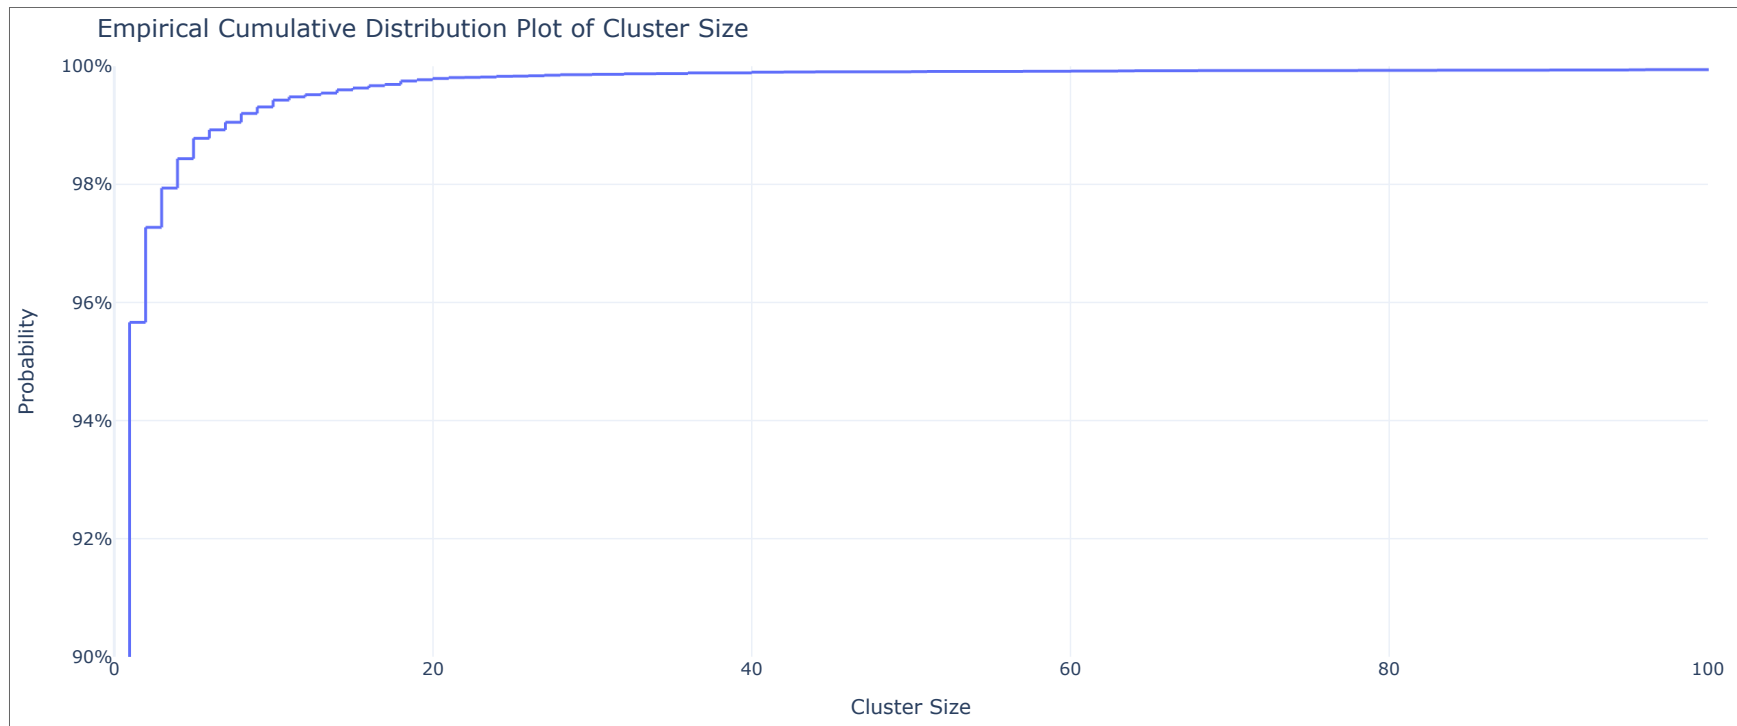
        if nearest_dist < MAX_DISTANCE:
            cluster_info.append((nearest_index, nearest_dist)) # Use nearest index as cluster
        else:
            cluster_info.append((index, 0.0)) # Create new cluster with current index
            forest.add(index, MinHash(hashvalues=sketch))

    return cluster_info
```

# Clustering: Results

In [429]:

```
cluster_size_ecdf_fig.show()
```



## Clustering: Extensions


1. **Partition clusters by creation time.** Clustering is an append only process. This opens up the opportunity to partition clusters by creation time to optimize database access.
2. **Detect and fix duplicate clusters.** If many similar errors get processed at the same time, they may end up creating many clusters. There is a need of a background job which will detect and fix duplicate clusters.
3. **Use different distance thresholds for different classes of errors.** Different frameworks may produce different logs and they may benefit from using criteria for near-duplicate failure.
4. **Use Least Recently Used (LRU) policy to manage cluster lifetime.** Most clusters represent unique errors and can be safely deleted after just a few days.

# Statistics

# Statistics: Top clusters by number of failed actions

In [430]:

```
top_clusters_by_number_of_actions_fig.show()
```



## Statistics: Top clusters by number of failed actions

In [367]:

```
print(actions.loc[335630, "action_output_failure"])
```

fatal: reference is not a tree: f5b4060fec2bd637a88bfafbcc16b87f0241656c

In [366]:

```
print(actions.loc[156330, "action_output_failure"])
```

```
$ jest --build -r=experimental --env=development --ci
```

- Unrecognized CLI Parameter:

Unrecognized option "build". Did you mean "bail"?

CLI Options Documentation:

<https://jestjs.io/docs/cli>

error Command failed with exit code 1.


info Visit <https://yarnpkg.com/en/docs/cli/run> for documentation about this command.



# Statistics: Top clusters by number of affected users

In [431]:

```
top_clusters_by_unique_users_fig.show()
```



## Statistics: Top clusters by number of affected users

Outage in external service:

In [378]:

```
print(actions.loc[295268, "action_output_failure"])
```

```
error An unexpected error occurred: "https://registry.yarnpkg.com/whatwg-fetch/-/whatwg-fetch-2.0.4.tgz: Request failed \"500 Internal Server Error\"".
info If you think this is a bug, please open a bug report with the information provided in "/home/circleci/project/yarn-error.log".
info Visit https://yarnpkg.com/en/docs/cli/install for documentation about this command.
```

Lint error in a commonly edited file(s):

In [384]:

```
print(actions.loc[309408, "action_output_failure"])
```

This project uses prettier to format all JavaScript code.

Please run `yarn prettier-all` and add changes to files listed below to your commit:

`packages/react-test-renderer/src/ReactTestRenderer.js`

# Integrations

1. **Incident Management** - widespread errors trigger incident management proc.
2. **Impact Assessment** - alerts have information about the number of affected users.
3. **Remediation Information** - Users can add remediation steps to error clusters.
4. **Error Suppression** - certain widespread errors can be suppressed and the the integration process can be resumed.
5. **Batch Retry** - once the root cause is resolved we can batch retry validation process for a precise set of blocked Code Reviews.
6. **Topline Metrics** - we estimate the impact of specific widespread errors on the topline metrics measuring developer productivity.

## Summary

1. Clustering of near-identical documents eliminates the need of maintaining a long list of normalization rules.
2. Minhash and LSH Forest are efficient solutions for detecting near-identical documents.  
Proposed settings proposed density of ~2M cluster per 1 GB.
3. For the React case study, during the 1 year period we observed 350k erros. Proposed pipeline produced 300k unique error clusters taking 150MB.
4. Proposed solution scales well and can be deployed in large CI systems.